Information Retrieval: From Classical Ranking Models to Neural and Conversational Search

Scientific overview paper

Abstract

Information retrieval (IR) studies how systems represent, index, search, rank, and evaluate access to information that is not neatly structured for direct machine use. The field underlies web search, enterprise search, legal retrieval, scientific discovery systems, recommendation-adjacent retrieval tasks, and increasingly retrieval-augmented generation. This paper surveys the foundations of IR, including indexing, relevance, ranking, and evaluation; explains the transition from lexical to neural approaches; and discusses recent developments in conversational search and retrieval for large language model pipelines. The core argument is that modern IR still rests on durable classical principles even as representation learning expands the field.

1. Introduction

Information retrieval is concerned with a deceptively simple question: how can a system return useful information for a human need expressed in incomplete, ambiguous, or underspecified language? Unlike structured database querying, retrieval assumes that both the information objects and the user's intent are uncertain. Documents may be long, noisy, and heterogeneous. Queries may be short, context-dependent, and vague. Relevance itself is not an intrinsic property of a document but a relationship between document, query, task, and user context. These characteristics distinguish IR from neighboring fields such as database systems, natural language processing, and recommendation, even though practical systems often combine elements of all three.

The web search engine is the most visible application of IR, but the field is broader. Enterprise search, digital libraries, legal research systems, biomedical literature platforms, e-commerce product search, and question answering pipelines all depend on retrieval. The objective is not simply to find exact matches, but to rank results so that the most useful items appear early. Since users rarely inspect more than a small fraction of the returned list, ranking quality matters at least as much as recall. This focus on ordered relevance explains why IR developed its own theories of term weighting, probabilistic ranking, and evaluation.

The field has changed substantially over the last decade. Representation learning, transformer models, dense retrieval, and cross-encoder rerankers have altered the technical landscape. At the same time, classical methods such as inverted indexes and BM25 remain central because they are efficient, interpretable, and robust. The practical architecture of many modern systems is hybrid: a lexical or sparse stage produces candidate documents, then a more expensive neural stage reranks them. Retrieval-augmented generation has further increased interest in IR because large language models depend heavily on the quality of retrieved context.

This paper reviews both the classical core and the contemporary expansion of the field. It begins with indexing and the probabilistic view of retrieval, moves to evaluation and benchmark culture, then discusses neural ranking, conversational search, and retrieval in generative pipelines. The overarching claim is that IR has not abandoned its foundations; it has layered new representation methods on top of them.

2. Indexing, Representation, and the Classical Retrieval Model

The practical success of retrieval systems begins with representation. Large collections are typically organized through inverted indexes that map terms to postings lists of documents or positions in which those terms appear. This structure enables fast candidate retrieval over corpora containing millions or billions of documents. Tokenization, normalization, stemming, stop-word handling, and field weighting all affect the final behavior of the system. Although these steps can appear mundane compared with neural modeling, they remain foundational because retrieval quality is strongly shaped by preprocessing assumptions.

The vector space model and related term-weighting schemes offered an early and influential account of relevance ranking. Documents and queries were represented as weighted term vectors, and similarity measures such as cosine similarity supported ranking. Term frequency-inverse document frequency captured an intuitive principle: terms are informative when they occur frequently in a document but not everywhere in the collection. This family of models created a workable bridge between linguistic intuition and efficient computation.

Probabilistic retrieval sharpened the picture by asking which document is most likely to be relevant given a query. The probabilistic relevance framework and later BM25 operationalized this idea in a form that proved both effective and durable. BM25 remains one of the strongest lexical baselines in modern retrieval and still anchors many production systems. Its persistence is not an accident. It handles term saturation and document length normalization sensibly, is computationally efficient, and often degrades gracefully when data are sparse or noisy.

Classical retrieval is sometimes caricatured as obsolete keyword matching. That caricature is false. Lexical systems encode important signals about exact term occurrence, collection statistics, and efficient candidate generation. Their weaknesses are real: synonymy, paraphrase, and semantic mismatch can cause relevant material to be missed. But these weaknesses are precisely why contemporary IR architectures often combine lexical precision with learned semantic representations rather than discard classical methods entirely.

3. Relevance, User Models, and Evaluation

Relevance is the conceptual center of IR, yet it is not a simple variable. Early formulations already distinguished topical relevance from situational usefulness, and later work emphasized graded relevance, novelty, diversity, and session context. A document may be topically related to a query but still unhelpful because it is redundant, outdated, unreadable, or aimed at a different task level. This complexity explains why evaluation metrics are task-dependent and why benchmark construction is intellectually significant rather than clerical.

Test collection methodology, especially through TREC, has been one of IR's most important institutional achievements. TREC created standardized corpora, topics, relevance judgments, and evaluation routines that made large-scale comparison possible. Official TREC overviews stress that the conference's function is not simply to declare winners but to build evaluation infrastructure for retrieval research. This infrastructure allowed steady progress in ad hoc retrieval, question answering, web search, conversational search, and more recent deep learning tracks.

Metrics such as precision, recall, mean average precision, normalized discounted cumulative gain, reciprocal rank, and recall at k operationalize different priorities. Precision-oriented tasks care about early ranking accuracy; recall-oriented tasks care about exhaustive coverage; graded metrics care about ordering among documents with different relevance levels. Choosing a metric is therefore a normative decision about what counts as system success. This becomes especially important when systems are deployed in medicine, law, or scientific search, where missing a crucial document may be more costly than placing a marginally relevant one too high.

User behavior complicates the picture further. Real users reformulate queries, inspect snippets, abandon sessions, and bring prior knowledge to the task. Conversational and interactive retrieval research emerged partly because one-shot query-document ranking is an incomplete account of search as practiced. Nevertheless, benchmark evaluation remains indispensable because without it the field would have little basis for comparing methods across groups. The tension between controlled evaluation and real-world interaction is therefore productive rather than fatal.

4. Neural Retrieval and Reranking

The neural turn in IR was driven by the desire to address vocabulary mismatch and to learn richer representations of relevance than term overlap alone could provide. Early neural ranking models learned dense vector representations for queries and documents, while later transformer-based approaches dramatically improved contextual matching. BERT-based rerankers, in particular, showed that deep contextual encoders could substantially outperform earlier neural models on benchmark datasets when applied to a candidate set generated by a first-stage retriever.

A practical distinction has emerged between bi-encoder dense retrieval and cross-encoder reranking. Dense retrievers encode queries and documents independently into vector spaces, enabling efficient approximate nearest-neighbor search. Cross-encoders jointly process query-document pairs and often achieve better ranking quality, but at much higher computational cost. This trade-off is why staged retrieval pipelines are now common. A sparse or dense retriever first narrows the search space; a stronger but slower model then reranks the candidates.

The rise of large training datasets accelerated this development. MS MARCO became especially influential because it offered large-scale relevance supervision for passage and document ranking. TREC’s deep learning tracks then provided benchmark settings in which different neural and hybrid methods could be compared rigorously. The official overviews of these tracks are clear on the point: the field was testing which methods work best in the large-training-data regime. Results consistently showed strong performance from architectures that combined robust first-stage retrieval with learned reranking.

Yet neural retrieval introduces its own problems. Dense models may be harder to interpret, more brittle under domain shift, and computationally expensive to update or deploy at scale. They can also amplify training-set biases. For these reasons, modern IR has not converged on a single dominant paradigm. Instead, it has converged on engineering pluralism: sparse, dense, and late-interaction methods coexist, and the best choice depends on latency budgets, collection characteristics, and task requirements.

5. Conversational Search and Retrieval-Augmented Generation

Conversational search extends retrieval from isolated queries to multi-turn interaction. Here the system must account for context, elliptical references, clarification needs, and evolving information goals. A user might ask a follow-up question that contains no explicit topic term, relying on the previous turn to supply the context. Classical ranking models are not designed for this directly, but multi-stage pipelines with query reformulation, context modeling, and reranking have produced workable solutions. TREC's conversational and related tracks helped transform these problems from conceptual curiosities into benchmarked research tasks.

Retrieval-augmented generation has created a new practical demand for IR. Large language models can produce fluent outputs, but without retrieval they are limited by static training data and may hallucinate unsupported claims. Retrieval augments generation by supplying relevant documents or passages at inference time. This reframes classic IR questions in a new environment: candidate recall, passage granularity, ranking diversity, latency, and evaluation all affect downstream answer quality. Recent TREC work on retrieval for RAG-like tasks shows that stronger retrievers often improve generated summaries, though metric choice can alter the conclusion.

These developments do not dissolve the distinction between IR and generation. Retrieval is still about selecting evidence; generation is about synthesizing or presenting it. Confusing the two leads to bad system design. A generative model cannot compensate reliably for poor evidence selection, and excellent retrieval does not automatically produce a faithful summary. The architecture must be evaluated end to end, but the retrieval component still deserves independent scrutiny.

Another issue is citation and provenance. In classical search, the user sees ranked documents directly. In RAG systems, the user may instead see an answer whose quality depends on partially hidden retrieval. This increases the importance of traceability, source attribution, and diagnostic evaluation. In a sense, modern generative systems have made classical IR concerns newly visible.

5A. Domain-Specific Retrieval and Knowledge Access

Retrieval systems behave differently across domains because corpora, user intents, and relevance criteria differ. In scientific and biomedical search, terminology is specialized, synonymy is pervasive, and the cost of omission can be high. Users may require comprehensive result sets rather than just attractive early precision. In legal retrieval, authority, citation structure, temporal validity, and jurisdiction influence relevance in ways that ordinary web benchmarks do not capture. Enterprise search introduces another set of constraints, including permissions, heterogeneous document formats, stale metadata, and organization-specific jargon. These domain conditions matter because a retrieval model that performs well on public web benchmarks may fail badly when the document distribution and task logic change.

Domain-specific retrieval often revives classical IR concerns in a more visible form. Metadata quality, controlled vocabularies, field weighting, and query expansion can outperform fashionable modeling choices when collections are narrow and carefully curated. At the same time, neural approaches may be especially helpful in domains characterized by paraphrase and terminology drift. The relevant lesson is not that one family of methods dominates, but that system design must begin with task analysis. Retrieval is highly sensitive to what counts as a useful miss, an acceptable false positive, and a credible explanation for ranking decisions.

Knowledge access systems increasingly combine search with recommendation, summarization, and interactive assistance. This creates both opportunity and risk. On the positive side, users can move more quickly from an information need to a usable synthesis. On the negative side, layers of ranking, aggregation, and generation can obscure where evidence came from and whether critical information was omitted. The more mediated the interface becomes, the more important explicit provenance and evaluation become.

For this reason, retrieval research remains strategically important in areas far beyond consumer search engines. Scientific discovery platforms, patent search, intelligence analysis, and compliance review all depend on retrieving the right evidence under uncertainty. Poor retrieval in these domains is not an inconvenience but a substantive epistemic failure.

5B. Efficiency, Indexing at Scale, and System Design

Large-scale retrieval is as much a systems problem as a modeling problem. Efficient indexing, compression, caching, query planning, and distributed execution determine whether a retrieval system can actually serve users under realistic latency budgets. Inverted indexes remain powerful because they support highly optimized term-at-a-time or document-at-a-time traversal strategies and can be distributed effectively. Approximate nearest-neighbor structures play an analogous role for dense retrieval, but they introduce their own tuning problems around recall, memory, and update cost. There is no universal winner; the correct architecture depends on collection size, update frequency, and acceptable trade-offs between quality and speed.

Freshness also complicates modern retrieval. Web and enterprise collections change continuously. Dense indexes and reranking caches may need costly updates, while lexical systems can often absorb change more easily. This is one reason hybrid retrieval has become attractive in practice. A sparse first stage offers efficient, update-friendly candidate generation, while a neural reranker adds semantic sensitivity only where it pays off most. The architecture reflects an engineering truth: perfect ranking quality is useless if the system cannot meet its latency or maintenance requirements.

Evaluation of efficiency must therefore accompany evaluation of effectiveness. A model that increases mean reciprocal rank slightly but multiplies compute cost and energy usage may be inferior in many deployment contexts. The broader field of IR has long understood this, but the recent enthusiasm for large neural models sometimes obscures it. Search infrastructure is not only about relevance metrics; it is about serviceable, debuggable, and updateable systems that deliver relevant results at scale.

This systems perspective also matters for multilingual and multimodal retrieval. Supporting diverse languages, OCR-heavy corpora, images, audio, or mixed media requires additional indexing and representation strategies. As search broadens beyond plain text, the classic retrieval agenda is not shrinking. It is becoming more general and more technically demanding.

5C. Open Problems

One unresolved problem is how to align retrieval evaluation with actual user benefit in settings where interaction is prolonged, conversational, or generation-mediated. Standard test collections remain indispensable, but they simplify context and often treat relevance as static. Better evaluation frameworks for sessions, explanation quality, provenance, and downstream decision support are still needed. Another open problem is robustness to adversarial manipulation, including search engine optimization spam, content farms, and automatically generated text designed to game ranking signals. As more content is machine-generated, distinguishing evidence quality from textual fluency will become even harder.

A further challenge concerns public trust. Search systems shape what users encounter first, which means ranking choices have epistemic and social consequences even when they are not explicitly political. Technical relevance models therefore interact with questions of transparency, accountability, and diversity of exposure. IR cannot solve all such questions by metric design alone, but it cannot ignore them either. The field’s next phase will need to combine strong engineering with more explicit reasoning about the conditions under which ranked access to information is trustworthy.

5D. Search Quality, Trust, and User Experience

Search quality is not exhausted by ranking metrics. Users judge systems through snippets, faceting options, response speed, result diversity, query suggestions, and the perceived trustworthiness of returned material. A retrieval system that surfaces useful documents but presents them opaquely may still fail its users.

Conversely, attractive interfaces can create a false sense of reliability if ranking errors are hidden behind polished presentation. The practical field of IR therefore intersects directly with user experience design.

Trust in retrieval is fragile because users rarely inspect the entire result set. They infer quality from the first few results and from how well the interface supports reformulation when initial attempts fail. This places a premium on explanation, provenance, and sensible query assistance. It also means that retrieval errors can have disproportionate effects on belief formation, particularly in domains where users treat ranking as an authority signal. From a systems perspective, high-quality IR should make it easier for users to understand why a result is present and what evidence underlies it.

As retrieval systems become embedded in assistants and generated-answer interfaces, preserving that trust will become harder. Users may no longer see the ranking process directly. For that reason, the classical IR insistence on evaluation, provenance, and error analysis is becoming newly relevant rather than less relevant. Search quality is, in the end, a problem of justified access to information.

6. Open Problems and Conclusion

Despite enormous progress, information retrieval still faces hard open problems. Relevance remains context-sensitive and often subjective. Domain adaptation is difficult because systems trained on web search behavior may not transfer cleanly to legal, biomedical, or multilingual settings. Evaluation is costly because relevance judgments are expensive and incomplete. Efficiency matters because state-of-the-art models can be computationally heavy. Fairness, transparency, and resistance to manipulation add further constraints in public-facing systems.

What is striking, however, is how stable the field's central ideas remain. Indexing, ranking, relevance, and evaluation still define the discipline. Neural methods expanded the representational tools available to the field, but they did not erase the need for strong baselines, benchmark culture, and careful error analysis. On the contrary, those classical practices are more necessary than ever because modern systems are more complex and opaque.

Information retrieval therefore remains one of the most practically significant areas of computer science. It mediates access to knowledge in science, commerce, law, medicine, and everyday life. Its technical questions are not merely about search boxes but about how people find and trust information in large collections. The field's enduring achievement is the disciplined treatment of relevance as a technical, empirical, and human problem at once.

Seen this way, IR is not a solved infrastructure layer sitting invisibly beneath newer AI systems. It is the discipline that determines whether large collections can be searched, compared, and trusted at all. As information environments become denser and more automated, that role only becomes more central.

References

Baeza-Yates, R., & Ribeiro-Neto, B. (2011). Modern Information Retrieval: The Concepts and Technology Behind Search (2nd ed.). Addison-Wesley.

Bajaj, P., Campos, D., Craswell, N., et al. (2016). MS MARCO: A human generated machine reading comprehension dataset. arXiv preprint arXiv:1611.09268.

Buckley, C., & Voorhees, E. M. (2005). Retrieval system evaluation. In TREC: Experiment and Evaluation in Information Retrieval (pp. 53-75). MIT Press.

Craswell, N., Mitra, B., Yilmaz, E., Campos, D., & Voorhees, E. M. (2020). Overview of the TREC 2020 deep learning track. In Proceedings of TREC 2020.

Craswell, N., Mitra, B., Yilmaz, E., Campos, D., & Voorhees, E. M. (2019). Overview of the TREC 2019 deep learning track. In Proceedings of TREC 2019.

Croft, W. B., Metzler, D., & Strohman, T. (2015). Search Engines: Information Retrieval in Practice. Pearson.

Manning, C. D., Raghavan, P., & Schutze, H. (2008). Introduction to Information Retrieval. Cambridge University Press.

NIST. (2024). Overview of TREC 2023. National Institute of Standards and Technology.

NIST. (2025). TREC overview and evaluation resources. National Institute of Standards and Technology.

Robertson, S. E., & Zaragoza, H. (2009). The probabilistic relevance framework: BM25 and beyond. Foundations and Trends in Information Retrieval, 3(4), 333-389.

Soboroff, I. (2022). Overview of TREC 2021. In Proceedings of TREC 2021.

Wang, J., et al. (2024). Utilizing BERT for information retrieval: Survey, applications, resources, and challenges. ACM Transactions on Information Systems, 42(4), 1-47.

Wang, Y., et al. (2024). Retrieval choices for summary generation in TREC 2024 retrieval-augmented generation tasks. In Proceedings of TREC 2024.

Zhai, C., & Lafferty, J. (2004). A study of smoothing methods for language models applied to information retrieval. ACM Transactions on Information Systems, 22(2), 179-214.

Zobel, J., & Moffat, A. (2006). Inverted files for text search engines. ACM Computing Surveys, 38(2), Article 6.